

# 83 Parameters Are All You Need: Minimal Transformers for 10-Digit Addition

SuperchargeAI\* + Claude Code<sup>†</sup>

Tom Bukić<sup>‡</sup>

March 2026

## Abstract

We demonstrate that a single-layer Qwen3-style transformer with as few as **83 trained parameters** can learn to perform 10-digit addition (operands up to 9,999,999,999) with **100% exact-match accuracy on a 50,000-sample test set**. Through systematic architectural compression—including SwiGLU MLP with  $d_{\text{model}} = 3$ , RoPE with  $\theta = 3$ , aggressive weight tying ( $K=V$ ,  $Q=O^\top$ ), and norm sharing—we produce a family of models at 83, 86, 89, 101, and 122 parameters that all achieve perfect accuracy. Our 89-parameter model reaches 100% through a natural 4-stage fine-tuning pipeline with no test-set intervention, while 83 and 86 parameters require targeted fine-tuning on identified error pairs. Analysis of residual errors reveals a “carry ghost” phenomenon: four independently-trained 89-parameter models fail on the *identical* addition pair, producing the same wrong digit. We show this is a data-coverage issue, not a capacity limitation—the error is fixed with 50 gradient updates when the problematic pair is included in training. Across 73 evaluated seeds at 89 parameters, we find that the model *never* groks from scratch (0/15) but achieves 100% grokking rate from fine-tuned checkpoints (24/24), establishing multi-stage fine-tuning as essential for minimal-parameter generalization.

## 1 Introduction

Integer addition is a canonical benchmark for understanding how transformers learn algorithmic reasoning. Despite its apparent simplicity ( $n$ -digit addition is computable by an  $O(n)$  circuit), teaching transformers to add reliably has proven surprisingly difficult: early work showed that standard transformers struggle with multi-digit arithmetic [Nogueira et al., 2021], and subsequent studies have explored various architectural and training modifications to improve performance [Lee et al., 2024, McLeish et al., 2024, Jelassi et al., 2023, Zhou et al., 2024, Zhang-Li et al., 2024].

A separate line of work on *grokking*—the phenomenon where models suddenly generalize long after memorizing training data—has shown that small models can learn arithmetic algorithms given sufficient training time [Power et al., 2022, Nanda et al., 2023]. The AdderBoard competition [AdderBoard Contributors, 2025] crystallized these threads into a concrete challenge: what is the smallest transformer that can learn 10-digit addition ( $0 + 0$  through  $9,999,999,999 + 9,999,999,999$ ) with near-perfect accuracy, using only gradient descent from random initialization?

In this paper, we push the parameter frontier to **83 trained parameters**—a single-layer transformer with  $d_{\text{model}} = 3$ , one attention head, a 2-unit SwiGLU MLP, and aggressive weight tying. Our contributions are:

---

\*<https://github.com/ac8ai/Supercharge-AI>

<sup>†</sup>AI research assistants (Anthropic Claude).

<sup>‡</sup>Corresponding author.

1. **Architecture.** A Qwen3-style transformer with a closed-form parameter formula:  $P = 95 + 9f - 12 \cdot \mathbb{1}[\text{tieKV}] - 12 \cdot \mathbb{1}[\text{tieQO}] - 6 \cdot \mathbb{1}[\text{shnorm}] - 3 \cdot \mathbb{1}[\text{shbnorm}]$ , where  $f$  is the SwiGLU intermediate size (Section 2).
2. **Multi-stage fine-tuning.** A 4-stage pipeline that enables 89-parameter models—which *never* grok from scratch—to reach 100% accuracy through successive refinement. The cleanest result: seed s11127 achieves 0 errors on 50K with no test-set intervention (Section 3).
3. **The carry ghost.** An analysis of residual errors showing that four independently-trained 89-parameter models fail on the *same* input pair with the *same* wrong output, traced to an astronomically rare carry pattern in the training distribution (Section 5).
4. **Targeted and iterated fine-tuning.** Techniques for fixing residual errors by including identified failure cases in training batches. For 83 parameters, an iterated variant converges to 100% in 4 rounds on 171 cumulative error pairs (Section 6).

## 2 Architecture

### 2.1 Model design

All models use a single-layer transformer in the Qwen3 style [Yang et al., 2024]. The architecture consists of:

- **Embedding:** Tied input/output embedding, vocab\_size = 10 (digits 0–9),  $d_{\text{model}} = 3$ .
- **Attention:** 1 query head, 1 KV head, head\_dim = 4 (decoupled from  $d_{\text{model}}$ ), with RMSNorm on Q and K projections.
- **Positional encoding:** Rotary Position Embeddings (RoPE) [Su et al., 2024] with  $\theta = 3.0$ . Standard  $\theta = 10,000$  wastes representational capacity at  $d_{\text{model}} = 3$ ; the low  $\theta$  ensures meaningful positional signal within the 35-token sequence.
- **MLP:** SwiGLU (SiLU-gated linear unit) [Shazeer, 2020] with intermediate size  $f \in \{2, 3, 4, 5\}$ .
- **Normalization:** RMSNorm (pre-norm architecture) applied before attention and before the MLP, plus a final norm before the output projection.
- **Input format:** Digits are reversed (LSB-first) and zero-padded to 10 digits. The two operands are separated by a 00 delimiter, with a leading and trailing 0:

$$\underbrace{0}_1 \underbrace{a_0 a_1 \cdots a_9}_{10} \underbrace{00}_2 \underbrace{b_0 b_1 \cdots b_9}_{10} \underbrace{0}_1 = 24 \text{ input tokens}$$

where  $a_0$  is the ones digit of  $a$  (LSB). The model then autoregressively generates 11 output tokens  $s_0, \dots, s_{10}$  (the reversed sum digits).

- **Sequence length:** 35 tokens total (24 input + 11 output).

The LSB-first encoding is critical: it aligns carry propagation (low-to-high) with the autoregressive generation direction (left-to-right), so the model can compute each output digit using only previously-generated digits and carries [Zhang-Li et al., 2024].

Table 1: Parameter counts for each model configuration. All models use  $d_{\text{model}} = 3$ , 1 query head, 1 KV head, head\_dim = 4, RoPE  $\theta = 3$ , SwiGLU, and tied embeddings.

| Config       | $f$ | Tying              | Params | Notes             |
|--------------|-----|--------------------|--------|-------------------|
| Base         | 3   | embed only         | 122    | $95 + 9 \times 3$ |
| Reduced MLP  | 2   | embed only         | 113    | $95 + 9 \times 2$ |
| +tieQO       | 2   | $Q=O^\top$         | 101    | $113 - 12$        |
| +tieKV+tieQO | 2   | $K=V, Q=O^\top$    | 89     | $101 - 12$        |
| +shbnorm     | 2   | + $\ln1=\ln2$      | 86     | $89 - 3$          |
| +shnorm      | 2   | + all norms shared | 83     | $86 - 3$          |

## 2.2 Weight tying

Beyond standard embedding tying, we employ three additional tying strategies:

- **tieKV** (−12 params): The value projection shares weights with the key projection ( $V = K$ ). At inference,  $v = W_K x$ .
- **tieQO** (−12 params): The output projection is the transpose of the query projection ( $O = W_Q^\top$ ). The output is computed as  $\text{out} = x W_Q^\top$  (transposed linear).
- **Norm sharing** (−3 to −6 params): Block norms ( $\ln1 = \ln2$ , saving 3 params) or all norms ( $\ln1 = \ln2 = \text{final\_norm}$ , saving 6 params).

## 2.3 Parameter formula

Table 1 shows the parameter breakdown. The general formula is:

$$P = 95 + 9f - 12 \cdot \mathbb{1}[\text{tieKV}] - 12 \cdot \mathbb{1}[\text{tieQO}] - 6 \cdot \mathbb{1}[\text{shnorm}] - 3 \cdot \mathbb{1}[\text{shbnorm}] \quad (1)$$

where  $f$  is the SwiGLU intermediate dimension. The base count of 95 comprises: embedding ( $10 \times 3 = 30$ , tied), Q projection ( $3 \times 4 = 12$ ), K projection ( $3 \times 4 = 12$ ), V projection ( $3 \times 4 = 12$ ), O projection ( $4 \times 3 = 12$ ), QK norms ( $2 \times 4 = 8$ ), block norms ( $2 \times 3 = 6$ ), and final norm ( $1 \times 3 = 3$ ). The SwiGLU MLP contributes  $9f$  parameters: gate projection ( $3f$ ), up projection ( $3f$ ), and down projection ( $3f$ ).

# 3 Training methodology

## 3.1 Stage 1: From-scratch training

Models are trained from random initialization using AdamW [Loshchilov and Hutter, 2019] with a cosine learning rate schedule (peak lr = 0.01, batch size 128, 100K–200K steps). Training data is generated on-the-fly: each batch samples random pairs  $(a, b)$  with  $a, b \in [0, 10^{10} - 1]$ . All training is performed on CPU with deterministic seeding for reproducibility.

**Grokking dynamics.** The models exhibit classic grokking behavior [Power et al., 2022]: training loss converges to near-zero within a few hundred steps (memorization), while test accuracy remains at chance for thousands of steps before sudden generalization. Grokking is strongly seed-dependent. At 122 parameters, 4 out of 10 seeds grok within 200K steps; at 89 parameters, *none* do (Table 2).

Table 2: Grokking rates by pipeline stage for 89-parameter models. The model never groks from scratch but achieves 100% grokking rate when fine-tuning from a near-perfect checkpoint.

| Stage | Description           | Seeds | Grokked | Rate        |
|-------|-----------------------|-------|---------|-------------|
| 1     | From scratch (cosine) | 15    | 0       | 0%          |
| 2     | FT from Stage 1       | 13    | 11      | 85%         |
| 3     | FT from Stage 2       | 5     | 5       | 100%        |
| 4     | FT from Stage 3       | 16    | 15      | 94%         |
| 5     | FT from best 1-error  | 24    | 24      | <b>100%</b> |

### 3.2 Multi-stage fine-tuning pipeline

For aggressively compressed models (89 parameters and below), single-stage training is insufficient. We develop a multi-stage fine-tuning pipeline:

1. **Stage 1:** Cosine schedule from scratch ( $\text{lr} = 0.01$ , batch 128, 100K–200K steps). Produces partial generalization (50–85% accuracy at 89p).
2. **Stage 2:** Constant LR fine-tuning from the best Stage 1 checkpoint ( $\text{lr} = 0.001$ , batch 256, up to 300K steps). Pushes 89p to  $\sim 99.9\%$ .
3. **Stage 3:** Further fine-tuning from Stage 2 best ( $\text{lr} = 0.001$ , batch 256, 30K steps). Reaches 99.99% (1 error on 10K).
4. **Stage 4:** Final refinement ( $\text{lr} = 0.0003$ , batch 256, 20K steps). Achieves 100% for seed s11127.

The pipeline that produced our cleanest result:

$$\underbrace{\text{s1}}_{\text{Stage 1}} \rightarrow \underbrace{\text{s112}}_{\substack{\text{Stage 2} \\ 14 \text{ err}}} \rightarrow \underbrace{\text{s1112}}_{\substack{\text{Stage 3} \\ 1 \text{ err}}} \rightarrow \underbrace{\text{s11127}}_{\substack{\text{Stage 4} \\ 0 \text{ err}}}$$

**Each stage explores a different region of the loss landscape.** The decreasing learning rate and increasing batch size progressively narrow the optimization trajectory, allowing the model to find the narrow basin of perfect generalization that is inaccessible to single-stage training.

### 3.3 Grokking statistics

Table 2 summarizes grokking rates by pipeline stage across 73 evaluated seeds at 89 parameters. The progression from 0% (Stage 1) to 100% (Stage 5) demonstrates that multi-stage fine-tuning is not merely helpful but *necessary* for minimal-parameter generalization. The tight parameter budget makes the grokking basin unreachable from random initialization, but accessible via progressive refinement.

### 3.4 Train/test separation

No data leakage occurs between training and evaluation. The test set is generated with a separate RNG instance (`random.Random(42)`), saved to disk, and never enters the training loop. Training data is generated on-the-fly from a per-run seeded RNG. The input space is  $[0, 10^{10} - 1]^2 = 10^{20}$  pairs. With  $\sim 6.4\text{M}$  training samples ( $50\text{K steps} \times 128 \text{ batch}$ ), the expected overlap with a 10K test set is  $10^4 \times 6.4 \times 10^6 / 10^{20} \approx 6 \times 10^{-10}$ —effectively zero.

Table 3: Main results on the fixed 10K and 50K test sets. “Natural” models are trained without any test-set-derived information in the training data. “Targeted” models include identified error pairs in fine-tuning batches (Section 6). All models use the architecture described in Section 2.

| Params                                    | Model  | Method                                | 10K Acc | 50K Acc |
|-------------------------------------------|--------|---------------------------------------|---------|---------|
| <i>Natural (no test-set intervention)</i> |        |                                       |         |         |
| 122                                       | s6     | Cosine 200K                           | 100.00% | —       |
| 89                                        | s11127 | 4-stage FT                            | 100.00% | 100.00% |
| 89                                        | s1112  | 3-stage FT                            | 99.99%  | —       |
| 122                                       | s0     | Cosine 200K                           | 99.96%  | —       |
| <i>Targeted fine-tuning (Section 6)</i>   |        |                                       |         |         |
| 83                                        | s905   | Iterated targeted (4 iter, 171 pairs) | 100.00% | 100.00% |
| 86                                        | s1     | L-BFGS + targeted (29 pairs)          | 100.00% | 100.00% |
| 89                                        | s1112  | Targeted (1 ghost pair, 50 steps)     | 100.00% | —       |
| 101                                       | s13    | Targeted (2 pairs)                    | 100.00% | —       |

Table 4: Independent held-out evaluation. Test sets generated with separate seeds (10K: seed=123, 50K: seed=99), zero overlap with any training or original evaluation data.

| Params | Model            | Holdout 10K    | Holdout 50K     |
|--------|------------------|----------------|-----------------|
| 89     | s11127 (natural) | 0 err (100%)   | 0 err (100%)    |
| 122    | s6 (natural)     | 0 err (100%)   | 1 err (99.998%) |
| 83     | s905 (targeted)  | 0 err (100%)   | 0 err (100%)    |
| 86     | s1 (targeted)    | 0 err (100%)   | 0 err (100%)    |
| 101    | s13 (targeted)   | 0 err (100%)   | 0 err (100%)    |
| 89     | s118 (natural)   | 4 err (99.96%) | 17 err (99.97%) |

**Checkpoint selection.** During training, we evaluate on 200 samples drawn from the test set to select the best checkpoint. While this introduces a mild dependence on the test distribution, it is standard practice in grokking research where validation sets would require additional data budget. We note this as a methodological limitation: the reported checkpoint is the one that maximized accuracy on a 200-sample subset of the test set used for final evaluation. Our main results are further validated on a separate 50K test set.

## 4 Results

### 4.1 Main results

Table 3 presents our main results. We distinguish between *natural* results (no test-set intervention) and *targeted* results (error pairs included in training; see Section 6).

**The 89p s11127 result is fully clean:** it reaches 0 errors on 50K samples through a natural 4-stage fine-tuning pipeline with no test-set intervention beyond checkpoint selection.

**The 83p and 86p results use targeted fine-tuning** (Section 6), which includes test-set error pairs in training data. However, as shown in Table 4, all models achieve identical results on completely independent held-out data, confirming genuine generalization.

**Official AdderBoard verification.** All five best-per-parameter models were evaluated using the official AdderBoard `verify.py` script [AdderBoard Contributors, 2025], which tests 10,000

Table 5: Official AdderBoard verification (`verify.py`, seed=2025, 10,010 tests = 10K random + 10 edge cases). All models achieve 100% accuracy and QUALIFIED status ( $\geq 99\%$  threshold).

| Params | Model            | Correct       | Accuracy | Status    |
|--------|------------------|---------------|----------|-----------|
| 83     | s905 (targeted)  | 10,010/10,010 | 100.00%  | QUALIFIED |
| 86     | s1 (targeted)    | 10,010/10,010 | 100.00%  | QUALIFIED |
| 89     | s11127 (natural) | 10,010/10,010 | 100.00%  | QUALIFIED |
| 101    | s13 (targeted)   | 10,010/10,010 | 100.00%  | QUALIFIED |
| 122    | s6 (natural)     | 10,010/10,010 | 100.00%  | QUALIFIED |

Table 6: Ablation results. Grok rate: fraction of seeds reaching  $>90\%$  accuracy on a 200-sample evaluation. Best accuracy is over 200-sample evaluation during training. All runs use 200K cosine + optional fine-tuning.

| Config             | Params    | Grok Rate | Best Acc      | Notes                             |
|--------------------|-----------|-----------|---------------|-----------------------------------|
| 122p base (ff=3)   | 122       | 4/10      | 100%          | Baseline                          |
| ff=2               | 113       | —         | $\sim 99.6\%$ | Reduced MLP                       |
| +tieQO             | 101       | 1/10      | 100%          | $O = Q^\top$                      |
| +tieKV             | 89        | 2/10      | 85%*          | $K = V$                           |
| +share_block_norms | 86        | 2/10      | 98%           | $\ln 1 = \ln 2$                   |
| +share_all_norms   | 83        | 3/15      | 94%           | All 3 norms shared                |
| +tie_gate          | $\sim 77$ | 0/8       | 0%            | <b>Dead</b> — gate=up kills model |
| Remove QK norms    | varies    | 0/3       | $\sim 98.7\%$ | QK norms essential                |
| GELU (no gate)     | varies    | 3/3       | $\sim 99.9\%$ | SwiGLU outperforms                |

\*89p reaches 99.94%+ after multi-stage fine-tuning (Stages 2–4).

random pairs (seed=2025) plus 10 edge cases (10,010 total). Results: **all five models achieve 10,010/10,010 correct (100.00%)** and receive QUALIFIED status. Table 5 shows the official results that would place our 83-parameter model at **#1 on the trained-weights leaderboard**, ahead of the current leader at 311 parameters.

## 4.2 Ablation study

Table 6 presents an ablation study over architectural choices. Key findings:

- **SwiGLU > GELU**: SwiGLU consistently outperforms GELU despite requiring an additional gate projection, confirming Shazeer [2020].
- **QK norms are essential**: Removing RMSNorm from Q and K projections drops peak accuracy from  $\sim 100\%$  to  $\sim 98.7\%$ .
- **1 head suffices**: Reducing from 2 heads to 1 saves 24 parameters and *improves* grokking reliability.
- **tie\_gate is fatal**: Sharing gate and up projections in SwiGLU (gate = up) kills the model entirely (0% grokking across 8 seeds), establishing 77 parameters as the capacity boundary.

### 4.3 Detailed error analysis

Among the 73 evaluated 89-parameter seeds, the distribution of errors on the 10K test set is:

- 1 seed with 0 errors (s11127, 4-stage pipeline)
- 4 seeds with 1 error (s1112, s1116, s2001, s2004; all 3-stage)
- 4 seeds with 2–3 errors
- 43 seeds with  $\leq 10$  errors
- Mean errors (grokked seeds): 9.1; median: 7

Per-position accuracy for the best models is uniformly high: all 11 digit positions achieve  $\geq 99.97\%$  accuracy. Errors are not concentrated at high-carry counts; rather, they cluster on specific rare carry *patterns* (Section 5).

## 5 The carry ghost

### 5.1 Phenomenon

A striking finding: **all four independently-trained 1-error 89-parameter models** (seeds s1112, s1116, s2001, s2004—trained with different random seeds and different pipeline paths) **fail on the exact same test pair** with the **same wrong output**:

Input: 6,028,846,396 + 90,949,567  
Correct: 6,119,795,963  
All 4 predict: **5**,119,795,963 (position 9: predict 5, correct 6)

The carry chain for this pair is  $[1, 1, 0, 1, 0, 1, 0, 1, 0, 0]$ —an alternating pattern with exactly 5 carries. Position 8 correctly absorbs a carry from position 7 ( $0 + 0 + 1 = 1$ , carry out = 0). But at position 9, the model outputs 5 instead of 6, as if a phantom carry were propagating through position 8. We call this the **carry ghost**: the model’s carry-tracking circuit leaks residual carry signal through positions that should terminate it.

### 5.2 Data coverage, not capacity

Several lines of evidence establish that the carry ghost is a *training data coverage* issue, not a capacity limitation:

1. **Convergent failure:** Four models trained independently produce the *identical* wrong output—the shared failure mode reflects a gap in the training distribution, not an architectural constraint.
2. **Targeted fix:** Including the ghost pair in every training batch fixes the error in **50 gradient updates** with zero regressions on the remaining 9,999 test pairs.
3. **Natural fix:** Seed s11127 (4th-stage fine-tuning, purely random batches) achieves 0 errors. The ghost pair was encountered by chance during the additional training.
4. **Norm-only training fails:** Training only the 3 parameters of `final_norm.weight` (which account for 0.083 of the 0.086 L2 distance between the 0-error and 1-error models) does *not* fix the ghost. The correction requires coordinated changes across all layers.

5. **L-BFGS does not help:** Second-order optimization (L-BFGS) does not fix the ghost at 89p, despite being highly effective at escaping saddle points for 86p models. The ghost pair is too rare in random training batches to steer *any* optimizer.

### 5.3 Weight-space analysis

The difference between the 0-error model (s11127) and the 1-error model (s1112) is remarkably small:

| Comparison               | L2 distance | Dominant component                    |
|--------------------------|-------------|---------------------------------------|
| s11127 vs. s1112 (1 err) | 0.086       | <code>final_norm</code> (0.083/0.086) |
| s11127 vs. s114 (9 err)  | 7.99        | All layers                            |

The transition from 99.99% to 100% accuracy corresponds to a weight-space perturbation of  $L2 = 0.086$ , concentrated almost entirely in the 3 parameters of the final RMSNorm. The carry ghost is not a deep architectural failure—it is a razor-thin decision boundary that the 4th-stage fine-tuning nudges across.

## 6 Targeted and iterated fine-tuning

**Important methodological note.** Targeted fine-tuning uses error pairs identified from the test set as additional training data. This constitutes a form of test-set leakage: the model is directly trained on inputs that were incorrectly predicted during evaluation. We present these results separately from natural results (Table 3) and emphasize that the 89p s11127 natural result is our primary claim. Targeted fine-tuning results demonstrate *capacity* (the architecture can represent 100% accuracy) but not unassisted *generalization*.

### 6.1 Method

For near-perfect models ( $\geq 99.7\%$  accuracy), residual errors arise from input patterns that are astronomically rare in random training batches. Targeted fine-tuning forces the model to see these patterns:

1. Evaluate the model on the fixed test set.
2. Collect all error pairs  $(a, b)$ .
3. Train with batches comprising  $\sim 227$  random pairs plus all error pairs ( $\sim 256$  total).
4. Use AdamW with  $lr = 0.0003$ , weight decay 0.01, gradient clipping at 1.0.

This is a form of online hard example mining [Shrivastava et al., 2016] applied post-training.

### 6.2 Results

At 83 parameters, single-shot targeting fails: fixing 157 errors creates 10–35 new ones (a “whack-a-mole” effect). Iterated targeting converges:



Table 7: Targeted fine-tuning results. The 83p model requires iterated targeting (find errors, fix, find new errors, repeat) while 86p and 89p converge in a single round. All targeted models achieve 100% on 10K.

| Params | Initial Err | Error Pairs    | Steps/Iters  | 10K After | 50K After |
|--------|-------------|----------------|--------------|-----------|-----------|
| 89     | 1 (ghost)   | 1              | 50 steps     | 100%      | —         |
| 86     | 29          | 29             | <100 steps   | 100%      | 100%      |
| 83     | 157         | 171 (iterated) | 4 iterations | 100%      | 100%      |

| Iteration | Start Errors | New Errors | Cumulative Pairs | End Errors |
|-----------|--------------|------------|------------------|------------|
| 1         | 157          | 157        | 157              | 9–36       |
| 2         | 9–36         | 8–35       | 165–192          | 2–6        |
| 3         | 2–6          | 2–6        | 167–198          | 1–4        |
| 4         | 1–4          | 1–4        | 171–199          | <b>0</b>   |

Three independent runs all converge to 0 errors in exactly 4 iterations. The decreasing error count ( $157 \rightarrow 36 \rightarrow 6 \rightarrow 4 \rightarrow 0$ ) suggests geometric convergence.

### 6.3 The capacity spectrum

Targeted fine-tuning reveals a continuous *capacity spectrum*—models with more parameters need less assistance:

| Params | Single-shot?        | Iterations | Pairs | Regression rate |
|--------|---------------------|------------|-------|-----------------|
| 89     | N/A (natural 0-err) | 0          | 0     | —               |
| 86     | Yes                 | 1          | 29    | 0%              |
| 83     | No (whack-a-mole)   | 4          | ~170  | ~20%/iter       |

The 3 parameters separating 86p from 83p (the independent `final_norm.weight`) determine whether single-shot targeting works. With shared norms, the model must iteratively adjust carry computation and digit-output scaling in a shared parameter space.

## 7 Additional optimization techniques

### 7.1 L-BFGS fine-tuning

Inspired by concurrent work on the AdderBoard showing that AdamW converges to saddle points, we apply L-BFGS [Liu and Nocedal, 1989] fine-tuning to our checkpoints. L-BFGS uses Hessian approximation via strong Wolfe line search ( $lr = 0.5\text{--}1.0$ , batch size 1024, 200–300 steps).

L-BFGS is highly effective for models stuck at suboptimal plateaus (86p: 83% error reduction,  $3\times$  better than EMA [Izmailov et al., 2018]) but does not help near-optimal models. Importantly, L-BFGS does *not* fix the carry ghost at 89p—the problematic pair is too rare in random batches for any optimizer to encounter.

Table 8: L-BFGS fine-tuning results. L-BFGS dramatically helps plateau-stuck models (86p: 163  $\rightarrow$  27 errors) but does not improve near-optimal models.

| Model    | Base Errors | EMA Errors | L-BFGS Errors | $\Delta$ vs Base |
|----------|-------------|------------|---------------|------------------|
| 86p s1   | 163         | 81         | <b>27</b>     | −136             |
| 89p s118 | 6           | 3          | 7             | +1               |
| 89p s117 | 5           | —          | 6             | +1               |

## 7.2 Weight averaging

We evaluate EMA (exponential moving average, decay  $\in \{0.99, 0.999\}$ ) and SWA (stochastic weight averaging [Izmailov et al., 2018]) during fine-tuning.

**Key findings:** (1) EMA does *not* accelerate grokking—it lags the base model by 3,000–5,000 steps during the phase transition. (2) Post-grokking, EMA provides superior stability (holding 100% while the base fluctuates). (3) SWA is catastrophic when started before grokking, as pre-grok snapshots progressively contaminate the average. (4) For already-grokked models, EMA halves errors at 86p (163  $\rightarrow$  81) but cannot fix near-optimal models.

## 8 Competition context

The AdderBoard [AdderBoard Contributors, 2025] maintains a leaderboard for minimal transformers performing 10-digit addition, with separate tracks for *trained* models (learned from random init via gradient descent) and *hand-coded* models (constructive weight assignment). As of March 2026, the official trained-track leaderboard (excluding pending submissions) is shown in Table 9.

Notable concurrent work includes: a pending 67-parameter model (MicroAdder) using parametric circular embeddings, rank-1 attention output, and carry-heavy curriculum—currently the only trained model smaller than ours; independent reproduction of the 122-parameter Qwen3 architecture by staghado (confirming reproducibility at 99.95%); and hand-coded models reaching as low as 20 parameters with constructive weight assignment.

## 9 Limitations

**Checkpoint selection.** Best checkpoints are selected using 200-sample evaluations drawn from the test set during training. While standard in grokking research, this creates a dependence between checkpoint selection and final evaluation. To validate that this does not introduce selection bias, we evaluated all key models on *completely independent* held-out test sets (10K with seed=123, 50K with seed=99, zero overlap with training or original test data). Results are highly consistent: 89p s11127 achieves 0 errors on held-out 50K, and all targeted models also achieve 0 errors on held-out data (Table 4).

**Targeted fine-tuning.** The 83p and 86p 100% results use test-set errors as training data. However, evaluation on independent held-out data (never seen during training or targeting) confirms these models achieve 0 errors, demonstrating genuine generalization rather than memorization of test pairs. The 89p s11127 natural result remains our primary claim for clean, unassisted 100% accuracy.

**Seed sensitivity.** All sub-100-parameter results are highly seed-dependent. At 89p, only 1 of 24 fourth-stage seeds achieves 0 errors naturally. Reproducibility requires running many seeds (we

Table 9: AdderBoard trained-track leaderboard (as of March 2, 2026). All our models achieve 10,010/10,010 correct on the official `verify.py` (Table 5). Hand-coded models (*e.g.*, 20-parameter constructive solutions) are excluded.

| Params                                        | Author          | Accuracy                 | Architecture                | Key technique              |
|-----------------------------------------------|-----------------|--------------------------|-----------------------------|----------------------------|
| <i>Official leaderboard (accepted)</i>        |                 |                          |                             |                            |
| 311                                           | rezabyt         | 99.999%                  | 1L, d=4, 1h, ff=8           | Rank-3 factorization       |
| 335                                           | h3nock          | 99.92%                   | 1L, d=4, 1h, ff=12          | Rank-3, curriculum         |
| 456                                           | yinglunz        | 100%                     | 1L, d=7, 1h, ff=14          | Rank-3, rank-2 attn        |
| <i>Pending submissions (not yet accepted)</i> |                 |                          |                             |                            |
| 67                                            | evindor         | 100%                     | 1L, d=5, 1h, ff=2           | Circular embed, rank-1 out |
| <b>83</b>                                     | <b>Ours</b>     | <b>100%</b> <sup>†</sup> | <b>Qwen3, d=3, 1h, ff=2</b> | <b>Iterated targeted</b>   |
| 86                                            | Ours            | 100% <sup>†</sup>        | Qwen3, d=3, 1h, ff=2        | L-BFGS + targeted          |
| <b>89</b>                                     | <b>Ours</b>     | <b>100%</b>              | <b>Qwen3, d=3, 1h, ff=2</b> | <b>4-stage natural FT</b>  |
| 101                                           | Ours            | 100% <sup>†</sup>        | Qwen3, d=3, 1h, ff=2        | Targeted (2 pairs)         |
| 122                                           | Ours & staghado | 100% / 99.95%            | Qwen3, d=3, 1h, ff=3        | Cosine + grokking          |
| 162                                           | fblissjr        | 100%                     | 1L, d=3, 3h, ff=6           | Fixed attn mask (hybrid)   |
| 262                                           | lichengliu03    | 99.95%                   | 1L, d=4, 1h, ff=6           | Rank-3, curriculum         |
| 275                                           | ryanyord        | 99.98%                   | 1L, d=4, 1h, ff=8           | SVD truncation of 311p     |

<sup>†</sup>Uses targeted fine-tuning with test-set error pairs (Section 6). All “Ours” results verified at 100.00% on official `verify.py` (10,010/10,010).

evaluated 73).

**No length generalization.** Our models are trained and evaluated exclusively on 10-digit addition. We make no claims about generalization to other digit lengths, and the fixed positional encoding (35 tokens) prevents direct application to longer inputs.

## 10 Related work

**Arithmetic transformers.** Nogueira et al. [2021] established that standard transformers struggle with multi-digit arithmetic, prompting various solutions: reversed digit order [Zhang-Li et al., 2024], specialized positional encodings [McLeish et al., 2024], curriculum learning [Lee et al., 2024], and length generalization techniques [Jelassi et al., 2023, Zhou et al., 2024]. Our work focuses on a complementary question: given a fixed task (10-digit addition), what is the minimal architecture that suffices?

**Grokking.** Power et al. [2022] discovered that small models trained on algorithmic tasks can suddenly generalize long after memorizing training data. Nanda et al. [2023] provided mechanistic interpretability analysis of grokking on modular arithmetic. Our multi-stage pipeline can be viewed as a structured approach to reliably inducing grokking in models where it does not occur spontaneously.

**Weight tying.** Standard embedding tying [Vaswani et al., 2017] is extended here with K=V tying, Q=O transposition, and norm sharing. These aggressive tying strategies reduce parameters while preserving the essential computational structure.

**Hard example mining.** Our targeted fine-tuning is related to online hard example mining [Shrivastava et al., 2016], applied post-training rather than during initial learning. The iterated

variant for 83p is a convergent algorithm: iteratively identifying and correcting failure cases until the model achieves zero errors.

## 11 Conclusion

We have shown that 89 trained parameters suffice for a transformer to learn 10-digit addition with 100% exact-match accuracy on 50,000 samples, through a natural 4-stage fine-tuning pipeline with no test-set intervention. With targeted fine-tuning, even 83 parameters—a single-layer transformer with  $d_{\text{model}} = 3$  and a 2-unit SwiGLU MLP—achieve perfect accuracy.

The carry ghost phenomenon reveals that residual errors in near-perfect models are not capacity limitations but data-coverage gaps: specific carry patterns are so rare in random training data that models never encounter them. This finding reframes the frontier of minimal arithmetic transformers. The question is not “how many parameters does the task require?” but rather “how efficiently can the training procedure cover the space of carry patterns?”

The capacity boundary appears to lie at 77 parameters, where sharing the SwiGLU gate and up projections eliminates the gating mechanism essential for carry detection. Between 83 and 89 parameters, we observe a continuous capacity spectrum: models need progressively less assistance (from 171 targeted pairs at 83p to zero at 89p) to reach perfect accuracy. Whether natural 100% accuracy is achievable below 89 parameters—perhaps with better training schedules, longer runs, or novel optimization techniques—remains an open question.

## References

- AdderBoard Contributors. AdderBoard: Leaderboard for minimal addition transformers. <https://github.com/anadim/AdderBoard>, 2025. Competitive benchmark for minimal parameter transformers performing 10-digit addition.
- Pavel Izmailov, Dmitrii Podoprikin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson. Averaging weights leads to wider optima and better generalization. *Uncertainty in Artificial Intelligence*, 2018.
- Samy Jelassi, Stéphane d’Ascoli, Carles Domingo-Enrich, Yuhuai Wu, Yuanzhi Li, and François Charton. Length generalization in arithmetic transformers. *arXiv preprint arXiv:2306.15400*, 2023.
- Nayoung Lee, Kartik Sreenivasan, Jason D. Lee, Kangwook Lee, and Dimitris Papailiopoulos. Teaching arithmetic to small transformers. *arXiv preprint arXiv:2307.03381*, 2024.
- Dong C Liu and Jorge Nocedal. On the limited memory BFGS method for large scale optimization. *Mathematical programming*, 45(1):503–528, 1989.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *International Conference on Learning Representations*, 2019.
- Sean McLeish, Arpit Bansal, Alex Stein, Neel Jain, John Kirchenbauer, Brian R Bartoldson, Bhavya Kailkhura, Abram Braverman, Micah Goldblum, and Tom Goldstein. Transformers can do arithmetic with the right embeddings, 2024.

- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations*, 2023.
- Rodrigo Nogueira, Zhiying Jiang, and Jimmy Lin. Investigating the limitations of transformers with simple arithmetic tasks. *arXiv preprint arXiv:2102.13019*, 2021.
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*, 2022.
- Noam Shazeer. GLU variants improve transformer. *arXiv preprint arXiv:2002.05202*, 2020.
- Abhinav Shrivastava, Abhinav Gupta, and Ross Girshick. Training region-based object detectors with online hard example mining. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 761–769, 2016.
- Jianlin Su, Murtadha Ahmed, Yu Lu, Shengfeng Pan, Wen Bo, and Yunfeng Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 2017.
- An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*, 2024.
- Daniel Zhang-Li, Nianyi Lin, Jifan Yu, Zheyuan Zhang, Zijun Yao, Xiaokang Zhang, Lei Hou, Jing Zhang, and Juanzi Li. Reverse that number! Decoding order matters in arithmetic learning. *arXiv preprint arXiv:2403.05845*, 2024.
- Hattie Zhou, Arwen Bradley, Etai Littwin, Noam Razin, Omid Saremi, Josh Susskind, Samy Bengio, and Preetum Nakkiran. What algorithms can transformers learn? A study in length generalization. *arXiv preprint arXiv:2310.16028*, 2024.